



Microsoft Fabric Tour

1



Title Sponsor



iLink Digital
Thinking Beyond

Gold Sponsor



iLink Digital

2



Direct to Insight: Unleashing Direct Lake Data Agents in MS Fabric

Chaudhary Ahmad Ali / Belinda Allen



3

Chaudhary Ahmad Ali



- Current working as Chief Architect at iLink Systems
- 24 plus years of experience in IT
 - 5 years as an SE, 10 years as a DE, and 10 years as a BI Engineer
- Recognitions & Certifications
 - Fast Track Recognized Solution Architect (FTRSA 2022-2023)
 - TOGAF Certified
 - PL-300, DP-500, DP-600, DP-700
- Favorite Movies
 - Meet Joe Black, Shawshank Redemption, Forrest Gump, Legends of the Fall
- Favorite Sports and Hobbies
 - Pickleball, Tennis, Billiard, spending time with family, traveling, and food





4

About Belinda Allen - @MsBelindaAllen

Belinda.Allen@iLink-Systems.com



From Alcolu, SC	• I lived the longest in New York City	
In Technology	• Director of Enablement – iLink Digital • Started in tech in 1991 • Microsoft Partner since 2001	
Microsoft MVP	• 13 years / 5 years GP and 8 years Power BI	
Microsoft MCT	• 13 years (with a hiatus in the middle)	
Co-Authored 2 books	• BI and Dynamics GP	
I love learning	• Technology, Wood Turning, Cooking	
I love hobbies	• Newest hobby: Pickleball • Favorite: Knitting/Crocheting	

 iLink Digital

5

Agenda



- Direct Lake
- Direct Query Vs Import Vs Direct Lake
- Power BI Development Framework (Direct Lake)
- Data Agents
- Direct Lake Best Practices
- Key factors to consider when selecting the appropriate storage mode
- Upcoming Direct Lake related features
- Q & A



6



Direct Lake

- **Loads only on demand data into memory and evicts unused data gradually**
- Another storage mode in Power BI, contains the features of the existing storage modes (Import and Direct Query)
- Can handle large data volumes with much better performance as compared to Direct Query.
- Saves the semantic model refresh time and compute cost
- Semantic Models & Reports availability time is reduced



7



Direct Lake Concepts

- Fabric (Or Premium) Capacity
- Parquet Files dependency
- V-Order
- Data Framing
- Data Transcoding
- Cold and Warm Cache
- Column Temperature
- Direct Query Fall Back
- Fixed Identity Recommended



8



Direct Lake Limitations

- Parquet Files per Delta Table
- Row Groups Per Delta Table
- Rows per table
- Max Memory (Including Warm Cache + Operational + Query Serving)
- No SQL Views
- RLS/OLS based delta tables are not supported (OneLake Security – future)
- No Excel Hierarchies & Drill throughs

Fabric/Power BI SKUs	Parquet files per table	Row groups per table	Rows per table (millions)	Max model size on disk/OneLake ¹ (GB)	Max memory (GB)
F2	1,000	1,000	300	10	3
F4	1,000	1,000	300	10	3
F8	1,000	1,000	300	10	3
F16	1,000	1,000	300	20	5
F32	1,000	1,000	300	40	10
F64/FT1/P1	5,000	5,000	1,500	Unlimited	25
F128/P2	5,000	5,000	3,000	Unlimited	50
F256/P3	5,000	5,000	6,000	Unlimited	100
F512/P4	10,000	10,000	12,000	Unlimited	200
F1024/P5	10,000	10,000	24,000	Unlimited	400
F2048	10,000	10,000	24,000	Unlimited	400

Direct Lake



9



Data Loading in Direct Lake and Import

- Direct Lake uses the VertiPaq engine, the same engine employed by Import mode, to load Delta table formats.
- The columnar compression logic and file structure of Delta tables are like the proprietary format in Import mode.
- Direct Lake involves data transcoding during the initial load into memory, which takes more time compared to Import mode, where data is directly loaded without transcoding.
- If data is evicted from memory in Direct Lake, it must be transcoded again, leading to slower reload times compared to Import mode. Additionally, data eviction is more aggressive in Direct Lake.
- When Direct Lake reaches its resource limits, it can fall back to Direct Query mode, whereas Import mode may encounter out-of-memory or resource errors.



10

DQ Vs Import Vs Direct Lake

Direct Query

Refers data in source

Pros

- Real Time Analysis
- Can handle large data volumes

Cons

- Report Rendering is Slow
- Dependent on data source optimization
- Limited support in complex DAX functions

Import

Loads all the data in memory upon Refresh

Pros

- Report Rendering is best
- Can handle complex calculations

Cons

- Data Latency
- Processing Complexity
- Processing Time

Direct Lake

Loads in use data in memory & gradually evicts unused data

Pros

- Real Time Analysis
- Can handle large data volumes
- Report rendering performance is fast.

Cons

- Row Volume Limits
- High cardinality columns initially loads slowly
- Works with Parquet only

Power BI Dev Framework (Direct Lake)

Analyze Data Sources

Data Profiling

Analyze Grain

Understand Calendar Dim

Prepare Bus Matrix

Data Retention

Data Refresh Policy

Load Data into LH/WH

Bronze Layer

- Bring Data in Original Format
- Append Data only
- Add metadata columns like Load Date and Load Id etc.

Silver Layer

- Data Cleansing
- Data Standardization
- Add Surrogate Keys
- Implement SCD if any

Gold Layer

- Implement Star Schema
- Apply Z-Order & V-Order

Semantic Model

Refer Gold Tables & create relations

Develop KPI & Hierarchies

RLS / OLS

Add Business friendly names, descriptions, & synonyms

Performance Optimization

Vertipaq Analyzer

Best Practice Analyzer

Semantic Link/Labs

Delta Analyzer

Performance Analyzer

Query Plans

Reports

Prepare Themes

Prepare Templates

Avoid Large Data Exports

Promote Analyze in Excel

Model Refresh

Trigger Smart Model Refresh (Reframing)

Pre Warm Cache

Send Notification

Fabric Data Agents



- **Talk to your Data in Onelake (Lakehouse, Warehouse, Eventhouse, Semantic Model)**
- **Conversation context** is preserved even in scenarios involving disparate data sources
- Create **Endpoints** which allows them to be consumed within Fabric or outside, with services CoPilot Studio, Azure AI Foundry and custom applications (coming soon).
- Existing security rules, such as Row-Level Security (RLS) and Column-Level Security (CLS) in OneLake sources, will be maintained and respected
- Can maintain, deploy and evaluate Data Agents programmatically using SDK
- Now Supported in all paid F-SKU (F2 and above)



13

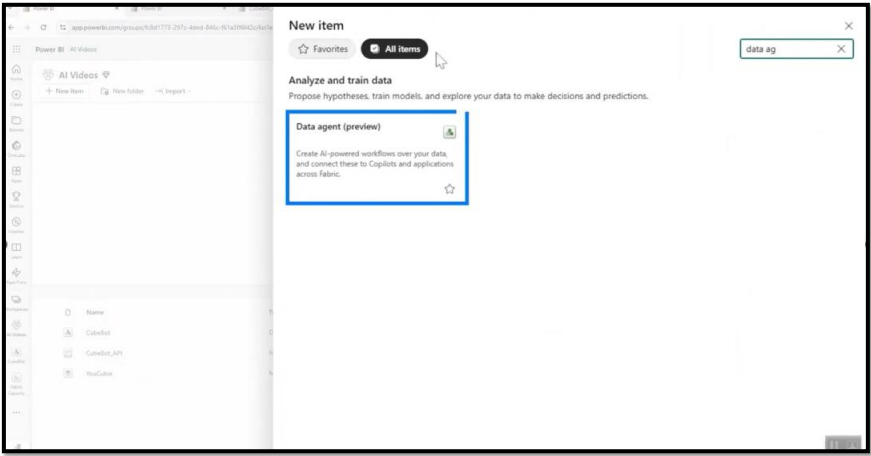
Fabric Data Agents (Semantic Model)



- **For Semantic Models type data sources, Adding instructions (up to 15,000 characters) is the most important configuration)**
- Natural Language to DAX (NL2DAX) is used to create DAX queries from conversations
- Establishing business-friendly table and column names, along with detailed descriptions and relevant synonyms, can significantly enhance the accuracy of DAX query generation.
- Granting access to Data Agents for other users also requires sharing semantic model-level access with the same users

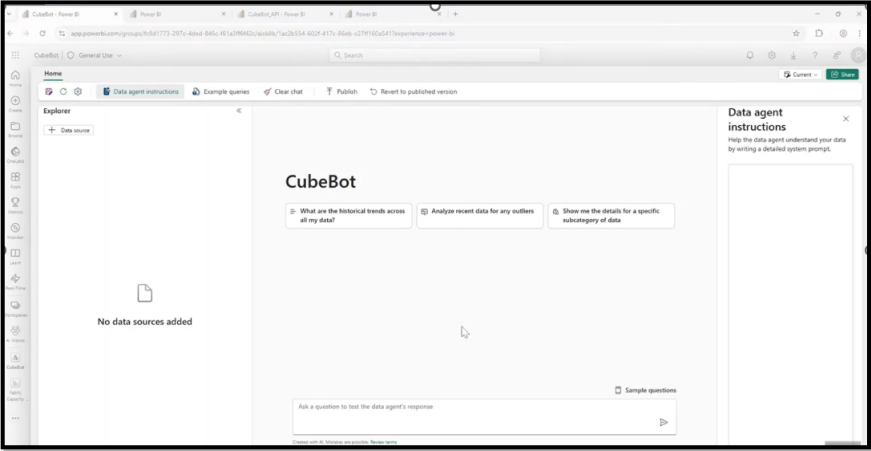


14



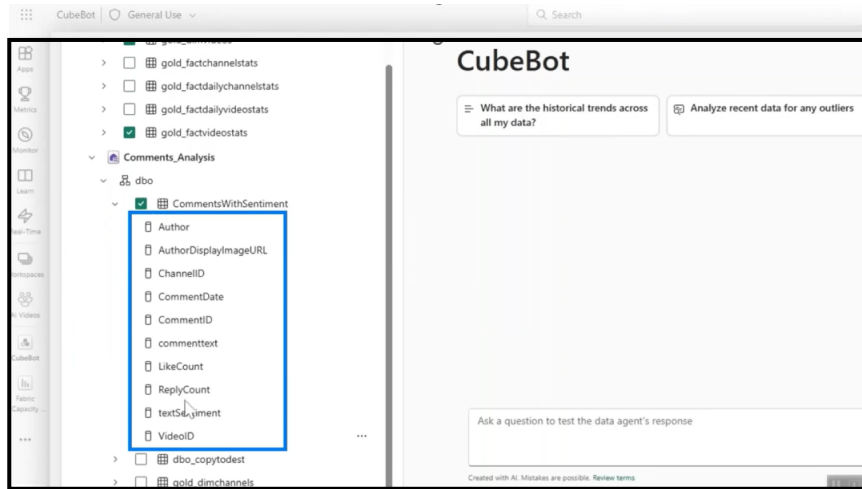
iLink Digital

15

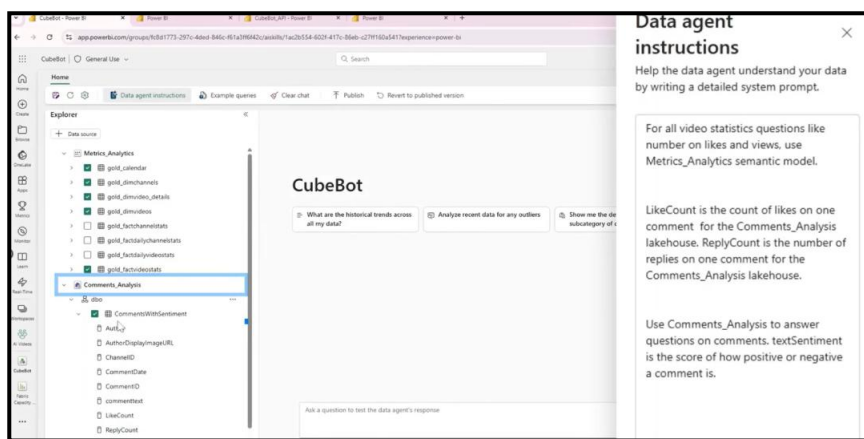


iLink Digital

16


 iLink Digital

17


 iLink Digital

18

Demo on Guy in a Cube (YouTube)



<https://www.youtube.com/watch?v=mFoSVcj6tLw>



19

Fabric Data Agents with Semantic Models



Semantic Model

Star Schema

Business Friendly Table,
Column and Measure Names

Define Descriptions

Provide Synonyms

Data Agents

Provide rich instructions

- Extract detail metadata from semantic model and can update the DA instructions using semantic Link Lab and DA SDK

Evaluate Data Agents
regularly

DA Integration

Within Fabric

Copilot Power BI

Azure AI Foundry

Teams

Copilot Studio (soon)



20

Direct Lake Best Practices (Gold Layer- LH)



- Implement Star Schema compliant tables
- Keep lean columns & rows, assign lean data type to columns
- Apply Z-Order and V-Order to fact tables to get best data storage and retrieval performance, e.g., Yellow Taxi Data 3 billion rows, CSV (**420 GB**) → Parquet (**168 GB**) → Delta with V-Order (**64 GB**)
- Parquet files from external sources such as Databricks or Snowflake are likely to perform slower when compared to V-Order optimized Parquet files in Fabric,
- Regular maintenance of gold layer delta tables with Optimize and Vacuum commands
- Configure higher bin size e.g. 500 MB to 1 GB in `spark.conf.set("spark.databricks.delta.optimizeWrite.binSize"` for partitioned tables



21

Direct Lake Best Practices (Semantic Model)



- Refer Star Schema delta tables in Gold Layer
- Avoid many to many and bidirectional relations
- Apply business user friendly table and column names. Consider defining synonyms, descriptions to make model CoPilot and Data Agent friendly.
- If high cardinality columns are occupying memory, consider following approaches:
 - Split columns e.g., datetime into data and time columns separately
 - Consider Precision points up to 3 point maximum in decimal type columns



23

Direct Lake Best Practices (Reports)



- Beware of Visual Rendering Time out setting (i.e. 225 seconds) on Fabric service is different than on Power BI desktop. Working visual in Power BI Desktop might fail in service in a case when data is first time loading/transcoding using direct lake and related column size is very high.
- Workaround, Pre-Warm the cache always using notebook scripts
- Leverage and Promote Data Agents for simple operational and analytical type questions over reports.



24

Direct Lake Best Practices (Security)



- Implement RLS or OLS in semantic model
- Using a Service Principal Name (SPN) to connect the Direct Lake semantic model to Lakehouse is highly recommended, particularly when Row-Level Security (RLS) is applied (One Lake Security will change this recommendation).



25

Direct Lake Best Practices (Refresh)



- Refreshing Direct Lake semantic model is also known as data reframing. Reframing is low-cost operation to reference latest version of data tables.
- Consider turning off auto direct lake refresh under semantic model settings.
- Trigger semantic model refresh using Fabric Data Factory or Notebook. Consider table-level refresh when multiple subject areas in the semantic model follow different data loading schedules.
- Identify most commonly used metric columns, Pre-Warm the cache for these columns. This will enhance the user experience even when they run the report for the first time.



26

Key Factors for selecting the Storage Mode



- Key Factors to consider before selecting the appropriate storage mode in Power BI:
 - Data Size (compressed and fact table row volumes)
 - Real Time Analytics required, justified by genuine business requirements
 - Report/Dashboards Rendering SLA
 - Report and Calculation Complexity (calculated tables or columns)
 - Data Source type
 - Is it an IT Driven Analytical Project?



27

Roadmap of Direct Lake and Data Agents



- Calculated Tables and Calculated Columns
- Add Delta tables from across the Lakehouse and across workspaces
- Implement Import and Direct Lake together in a same model
- Materialized Views support in Lakehouse
- One Lake Security
- Data Agents Integration with Copilot Studio
- More Sources will be supported with Data Agents like unstructured data sources or mirrored data sources.
- Org Apps support for the Data Agents
- (More)



28

Conclusion



- Direct Lake is another storage mode, which cannot replace existing storage modes.
- Consider Import Mode by default unless capacity memory limits are reached
- Consider Direct Lake when columns can fit within memory limits, while Import mode exceeds capacity constraints. Regularly track memory usage of Direct Lake models through DMV storage views.
- Choose Direct Query when the model exceeds memory limits and real-time analytics are essential.
- Direct Query is the slowest in terms of performance among the three options. Direct Lake performs nearly on par with Import mode, but Import consistently delivers the best performance overall.



29



Q & A & Thank You



iLink Digital
Thinking Beyond



30

Presentation Link



31

More Enablement options
<https://forms.office.com/r/UDRZaBHyz6>



32